

# Documentation for the Multivariate Normal Gibbs sampler

Anand

February 19, 2007

The Gibbs samplers in `Normal.py` can be applied to parameters  $x$  whose children are distributed as follows: For each child  $c_i$ ,

$$(c_i|x, F_i, a_i, \tau_i) \sim N(F_i x + a_i, \tau_i).$$

## 1 Conjugate

The sampling method `cmvNormalWithMVNormalChildren` applies if  $x$ 's prior is normal:

$$(x|\mu_p, \tau_p) \sim N(\mu_p, \tau_p).$$

In this case, the log conditional probability up to constants not involving  $x$ ,  $\sum_i \log p_i(c_i|x, \dots) + \log p_x(x|\dots)$ , is equal to

$$-\frac{1}{2}(c_i - F_i x - a_i)^T \tau_i (c_i - F_i x - a_i) - \frac{1}{2}(x - \mu_p)^T \tau_p (x - \mu_p),$$

which after completing the square yields

$$\begin{aligned} (x|\{c_i\}, \dots) &\sim N(\mu, \tau), \\ \tau &= \tau_p + \sum_i F_i^T \tau_i F_i, \\ \mu &= \tau^{-1}(\tau_p \mu_p + \sum_i F_i^T \tau_i (x_i - a_i)). \end{aligned}$$

This sampling method's constructor arguments are:

- **parameter**: The parameter representing  $x$ .
- **F\_dict**: A dictionary, indexed by child objects, whose values are nodes whose values are the  $F$  matrices.
- **a\_dict**: Same as **F\_dict**, but the values of the nodes are the  $a$  arrays.
- **tau\_dict**: The obvious.
- **prior\_mu**: A node whose value is the prior mean of  $x$ .
- **prior\_tau**: Same as **prior\_mu**, but with the precision matrix.

## 2 Nonconjugate

The sampling method `cmvNormalWithMVNormalChildren` applies if  $x$ 's prior  $p_x(x)$  is non-normal. In this case, the log joint probability of  $x$  up to unimportant constants,  $\sum_i \log p_i(c_i|x, \dots) + \log p_x(x)$ , is equal to

$$-\sum_i \frac{1}{2} (c_i - F_i x - a_i)^T \tau_i (c_i - F_i x - a_i) + \log p_x(x).$$

A Metropolis-Hastings step algorithm for this distribution which tends to be efficient is the following:

1. Propose a value  $x^p$  for  $x$  as if it had an uninformative prior ( $\tau_p = \epsilon I$ , with  $\epsilon \ll 1$ ). That is,

$$\begin{aligned} (x_p | \{c_i\}, \dots) &\sim N(\mu, \tau), \\ \tau &= \sum_i F_i^T \tau_i F_i, \\ \mu &= \tau^{-1} \sum_i F_i^T \tau_i (x_i - a_i). \end{aligned}$$

2. Accept the jump with probability

$$\min \left\{ 1, \frac{p_x(x^p)}{p_x(x)} \right\}.$$

This sampling method's constructor arguments are simply:

- **parameter**: The parameter representing  $x$ .
- **F\_dict**: A dictionary, indexed by child objects, whose values are nodes whose values are the  $F$  matrices.
- **a\_dict**: Same as **F\_dict**, but the values of the nodes are the  $a$  arrays.
- **tau\_dict**: The obvious.

## 3 Things that would be nice for Gibbs sampling in general

I started a class called `Gibbs` in `PyMC2/special_SamplingMethods/GibbsSampler.py`, but couldn't figure out how to make it work well. It would be nice if:

- We come up with some consistent class factory for making conjugate/ nonconjugate pairs of Gibbs samplers, maybe similar to what we're doing in `distributions.py`.
- The accessory dictionaries that get passed in to the constructor should be allowed to contain parameters, nodes, or simple ndarrays. The PyMC objects should have their values extracted automatically somehow.
- Are we willing to insist on consistent parent keying schemes, to avoid having to pass in the accessory dictionaries? If so, we need to standardize across all Gibbs samplers and `distributions.py`. Compatibility for weird models will require heavy use of utility nodes.
- Each Gibbs sampler holds a `OneAtATimeMetropolis` instance called 'default' or something which it uses if a Gibbs step fails.